

Import libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Read the Dataset

```
In [3]: df=pd.read_csv('heart.csv')
```

```
In [5]: df
```

```
Out[5]:
```

	Unnamed: 0	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exersa
0	0	0.479167	1.0	0.000000	0.292453	0.196347	0.0	0.5	0.740458	0.0
1	1	0.500000	1.0	0.000000	0.433962	0.175799	1.0	0.0	0.641221	0.0
2	2	0.854167	1.0	0.000000	0.481132	0.109589	0.0	0.5	0.412214	0.0
3	3	0.666667	1.0	0.000000	0.509434	0.175799	0.0	0.5	0.687023	0.0
4	4	0.687500	0.0	0.000000	0.415094	0.383562	1.0	0.5	0.267176	0.0
...
1020	1020	0.625000	1.0	0.333333	0.433962	0.216895	0.0	0.5	0.709924	0.0
1021	1021	0.645833	1.0	0.000000	0.292453	0.301370	0.0	0.0	0.534351	0.0
1022	1022	0.375000	1.0	0.000000	0.150943	0.340183	0.0	0.0	0.358779	0.0
1023	1023	0.437500	0.0	0.000000	0.150943	0.292237	0.0	0.0	0.671756	0.0
1024	1024	0.520833	1.0	0.000000	0.245283	0.141553	0.0	0.5	0.320611	0.0

1025 rows × 15 columns



```
In [6]: df.shape
```

```
Out[6]: (1025, 15)
```

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Unnamed: 0   1025 non-null   int64
 1   age          1025 non-null   float64
 2   sex          1025 non-null   float64
 3   cp           1025 non-null   float64
 4   trestbps     1025 non-null   float64
 5   chol         1025 non-null   float64
 6   fbs          1025 non-null   float64
 7   restecg      1025 non-null   float64
 8   thalach      1025 non-null   float64
 9   exang        1025 non-null   float64
10   oldpeak      1025 non-null   float64
11   slope        1025 non-null   float64
12   ca           1025 non-null   float64
13   thal         1025 non-null   float64
14   target       1025 non-null   float64
dtypes: float64(14), int64(1)
memory usage: 120.2 KB
```

```
In [8]: df.columns
```

```
Out[8]: Index(['Unnamed: 0', 'age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg',
              'thalach', 'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],
              dtype='object')
```

dropping unwanted columns

```
In [10]: df=df.drop(columns=['ca','fbs','restecg','Unnamed: 0'],axis=1)
```

```
In [11]: df.columns
```

```
Out[11]: Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'thalach', 'exang', 'oldpeak',
              'slope', 'thal', 'target'],
              dtype='object')
```

```
In [12]: df.describe()
```

Out[12]:

	age	sex	cp	trestbps	chol	thalach	
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1
mean	0.529878	0.695610	0.314146	0.354827	0.273973	0.596291	
std	0.189006	0.460373	0.343214	0.165252	0.117791	0.175616	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.395833	0.000000	0.000000	0.245283	0.194064	0.465649	
50%	0.562500	1.000000	0.333333	0.339623	0.260274	0.618321	
75%	0.666667	1.000000	0.666667	0.433962	0.340183	0.725191	
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	

In [13]: `df.duplicated().sum()`

Out[13]: 723

In [14]: `df_unique = df.drop_duplicates()`In [15]: `df_unique`

Out[15]:

	age	sex	cp	trestbps	chol	thalach	exang	oldpeak	slope	
0	0.479167	1.0	0.000000	0.292453	0.196347	0.740458	0.0	0.161290	1.0	1.0
1	0.500000	1.0	0.000000	0.433962	0.175799	0.641221	1.0	0.500000	0.0	1.0
2	0.854167	1.0	0.000000	0.481132	0.109589	0.412214	1.0	0.419355	0.0	1.0
3	0.666667	1.0	0.000000	0.509434	0.175799	0.687023	0.0	0.000000	1.0	1.0
4	0.687500	0.0	0.000000	0.415094	0.383562	0.267176	0.0	0.306452	0.5	0.6
...
723	0.812500	0.0	0.666667	0.245283	0.194064	0.335878	0.0	0.241935	0.5	0.6
733	0.312500	0.0	0.666667	0.132075	0.034247	0.793893	0.0	0.096774	0.5	0.6
739	0.479167	1.0	0.000000	0.320755	0.294521	0.687023	1.0	0.000000	1.0	1.0
843	0.625000	1.0	1.000000	0.622642	0.335616	0.412214	0.0	0.000000	1.0	0.6
878	0.520833	1.0	0.000000	0.245283	0.141553	0.320611	0.0	0.225806	0.5	1.0

302 rows × 11 columns

In [16]: `df.isnull().sum()`

```
Out[16]: age      0
sex      0
cp       0
trestbps 0
chol     0
thalach  0
exang    0
oldpeak  0
slope    0
thal     0
target   0
dtype: int64
```

MinMax Normalization

```
In [17]: from sklearn.preprocessing import MinMaxScaler
```

```
In [18]: scaler = MinMaxScaler()
df_scaled = pd.DataFrame(scaler.fit_transform(df), columns=df.columns)
df_scaled
```

```
Out[18]:
```

	age	sex	cp	trestbps	chol	thalach	exang	oldpeak	slope	
0	0.479167	1.0	0.000000	0.292453	0.196347	0.740458	0.0	0.161290	1.0	1.
1	0.500000	1.0	0.000000	0.433962	0.175799	0.641221	1.0	0.500000	0.0	1.
2	0.854167	1.0	0.000000	0.481132	0.109589	0.412214	1.0	0.419355	0.0	1.
3	0.666667	1.0	0.000000	0.509434	0.175799	0.687023	0.0	0.000000	1.0	1.
4	0.687500	0.0	0.000000	0.415094	0.383562	0.267176	0.0	0.306452	0.5	0.
...
1020	0.625000	1.0	0.333333	0.433962	0.216895	0.709924	1.0	0.000000	1.0	0.
1021	0.645833	1.0	0.000000	0.292453	0.301370	0.534351	1.0	0.451613	0.5	1.
1022	0.375000	1.0	0.000000	0.150943	0.340183	0.358779	1.0	0.161290	0.5	0.
1023	0.437500	0.0	0.000000	0.150943	0.292237	0.671756	0.0	0.000000	1.0	0.
1024	0.520833	1.0	0.000000	0.245283	0.141553	0.320611	0.0	0.225806	0.5	1.

1025 rows × 11 columns



Statistical feature extraction

```
In [19]: features = df.drop('target', axis=1)
target_variables = df['target']

# Statistical feature extraction (row-wise)
stat_features = pd.DataFrame()
stat_features['mean'] = features.mean(axis=1)
stat_features['std'] = features.std(axis=1)
```

```

stat_features['min'] = features.min(axis=1)
stat_features['max'] = features.max(axis=1)
stat_features['range'] = features.max(axis=1) - features.min(axis=1)
stat_features['median'] = features.median(axis=1)

# Quantiles (25th and 75th)
stat_features['25%'] = features.quantile(0.25, axis=1)
stat_features['75%'] = features.quantile(0.75, axis=1)

# Variance
stat_features['variance'] = features.var(axis=1)

stat_features

```

Out[19]:

	mean	std	min	max	range	median	25%	75%	variance
0	0.486971	0.415874	0.0	1.0000	1.0000	0.385810	0.170054	0.935115	0.172951
1	0.525098	0.390599	0.0	1.0000	1.0000	0.500000	0.240340	0.910305	0.152568
2	0.527646	0.412954	0.0	1.0000	1.0000	0.450243	0.185245	0.963542	0.170531
3	0.503892	0.430478	0.0	1.0000	1.0000	0.588050	0.043950	0.921756	0.185311
4	0.322645	0.260517	0.0	0.6875	0.6875	0.345007	0.066794	0.478774	0.067869
...
1020	0.598578	0.350013	0.0	1.0000	1.0000	0.645833	0.358491	0.927481	0.122509
1021	0.572562	0.342323	0.0	1.0000	1.0000	0.517176	0.338931	0.911458	0.117185
1022	0.455286	0.342452	0.0	1.0000	1.0000	0.366889	0.206013	0.625000	0.117273
1023	0.321910	0.358861	0.0	1.0000	1.0000	0.221590	0.000000	0.609375	0.128781
1024	0.395409	0.363894	0.0	1.0000	1.0000	0.282947	0.162616	0.515625	0.132419

1025 rows × 9 columns

Concatenate normalized data and statistical features

```

In [20]: # Concatenate normalized data and statistical features
combined_output = pd.concat([df_scaled, stat_features], axis=1)

print(combined_output.head())

```

	age	sex	cp	trestbps	chol	thalach	exang	oldpeak	slope	\
0	0.479167	1.0	0.0	0.292453	0.196347	0.740458	0.0	0.161290	1.0	
1	0.500000	1.0	0.0	0.433962	0.175799	0.641221	1.0	0.500000	0.0	
2	0.854167	1.0	0.0	0.481132	0.109589	0.412214	1.0	0.419355	0.0	
3	0.666667	1.0	0.0	0.509434	0.175799	0.687023	0.0	0.000000	1.0	
4	0.687500	0.0	0.0	0.415094	0.383562	0.267176	0.0	0.306452	0.5	

	thal	target	mean	std	min	max	range	median	\
0	1.000000	0.0	0.486971	0.415874	0.0	1.0000	1.0000	0.385810	
1	1.000000	0.0	0.525098	0.390599	0.0	1.0000	1.0000	0.500000	
2	1.000000	0.0	0.527646	0.412954	0.0	1.0000	1.0000	0.450243	
3	1.000000	0.0	0.503892	0.430478	0.0	1.0000	1.0000	0.588050	
4	0.666667	0.0	0.322645	0.260517	0.0	0.6875	0.6875	0.345007	

	25%	75%	variance
0	0.170054	0.935115	0.172951
1	0.240340	0.910305	0.152568
2	0.185245	0.963542	0.170531
3	0.043950	0.921756	0.185311
4	0.066794	0.478774	0.067869

splitting Training and testing data

```
In [21]: from sklearn.model_selection import train_test_split
X=combined_output.drop('target',axis=1)
y=combined_output['target']

# First split
X_train1, X_test1, y_train1, y_test1 = train_test_split(X, y, test_size=0.2, ran

# Second split
X_train2, X_test2, y_train2, y_test2 = train_test_split(X, y, test_size=0.3, ran
```

```
In [22]: from sklearn.tree import DecisionTreeClassifier

from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (accuracy_score, precision_score, recall_score, conf
```

```
In [23]: models1=[('RandomForest',RandomForestClassifier()),
                  ('LogisticRegression',LogisticRegression()),
                  ('DecisionTree',DecisionTreeClassifier())]
```

```
In [24]: models2=[('RandomForest',RandomForestClassifier()),
                  ('LogisticRegression',LogisticRegression()),
                  ('DecisionTree',DecisionTreeClassifier())]
```

Define functions for additional metrics

```
In [25]: # Define functions for additional metrics
def specificity_score(conf_matrix):
    tn, fp, fn, tp = conf_matrix.ravel()
    return tn / (tn + fp)

def npv_score(conf_matrix):
```

```

    tn, fp, fn, tp = conf_matrix.ravel()
    return tn / (tn + fn)

def fpr_score(conf_matrix):
    tn, fp, fn, tp = conf_matrix.ravel()
    return fp / (fp + tn)

def fnr_score(conf_matrix):
    tn, fp, fn, tp = conf_matrix.ravel()
    return fn / (fn + tp)

def fmeasure_score(precision, recall):
    return 2 * (precision * recall) / (precision + recall)

# Function to evaluate models
def evaluate_models(models, X_train, X_test, y_train, y_test):
    for name, model in models:
        print(f"{name}")
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)

        conf_matrix = confusion_matrix(y_test, y_pred)

        print(conf_matrix)
        print("Accuracy:", accuracy_score(y_test, y_pred))
        print("Precision:", precision_score(y_test, y_pred))
        print("Recall:", recall_score(y_test, y_pred))
        print("Fmeasure:", fmeasure_score(precision_score(y_test, y_pred), recall_score(y_test, y_pred)))
        print("Sensitivity:", recall_score(y_test, y_pred))
        print("Specificity:", specificity_score(conf_matrix))
        print("MCC:", matthews_corrcoef(y_test, y_pred))
        print("NPV:", npv_score(conf_matrix))
        print("FPR:", fpr_score(conf_matrix))
        print("FNR:", fnr_score(conf_matrix))
        print('\n')

# Evaluate models on both splits
print("First split(80-20):")
evaluate_models(models1, X_train1, X_test1, y_train1, y_test1)

print(" Second split(70-30):")
evaluate_models(models2, X_train2, X_test2, y_train2, y_test2)

```

First split(80-20):

RandomForest

[[98 0]

[0 107]]

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

Fmeasure: 1.0

Sensitivity: 1.0

Specificity: 1.0

MCC: 1.0

NPV: 1.0

FPR: 0.0

FNR: 0.0

LogisticRegression

[[81 17]

[12 95]]

Accuracy: 0.8585365853658536

Precision: 0.8482142857142857

Recall: 0.8878504672897196

Fmeasure: 0.8675799086757991

Sensitivity: 0.8878504672897196

Specificity: 0.826530612244898

MCC: 0.7167775340329366

NPV: 0.8709677419354839

FPR: 0.17346938775510204

FNR: 0.11214953271028037

DecisionTree

[[98 0]

[0 107]]

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

Fmeasure: 1.0

Sensitivity: 1.0

Specificity: 1.0

MCC: 1.0

NPV: 1.0

FPR: 0.0

FNR: 0.0

Second split(70-30):

RandomForest

[[145 0]

[6 157]]

Accuracy: 0.9805194805194806

Precision: 1.0

Recall: 0.9631901840490797

Fmeasure: 0.98125

Sensitivity: 0.9631901840490797

Specificity: 1.0

MCC: 0.9617264301269595

NPV: 0.9602649006622517

FPR: 0.0

FNR: 0.03680981595092025


```

LogisticRegression
[[119  26]
 [ 20 143]]
Accuracy: 0.8506493506493507
Precision: 0.8461538461538461
Recall: 0.8773006134969326
Fmeasure: 0.8614457831325302
Sensitivity: 0.8773006134969326
Specificity: 0.8206896551724138
MCC: 0.700126342832296
NPV: 0.8561151079136691
FPR: 0.1793103448275862
FNR: 0.12269938650306748

```

```

DecisionTree
[[145  0]
 [ 6 157]]
Accuracy: 0.9805194805194806
Precision: 1.0
Recall: 0.9631901840490797
Fmeasure: 0.98125
Sensitivity: 0.9631901840490797
Specificity: 1.0
MCC: 0.9617264301269595
NPV: 0.9602649006622517
FPR: 0.0
FNR: 0.03680981595092025

```

Function to collect all metrics

```

In [26]: # Function to collect all metrics
def get_all_metrics(models, X_train, X_test, y_train, y_test, split_name):
    metrics = []
    for name, model in models:
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)

        conf_matrix = confusion_matrix(y_test, y_pred)
        accuracy = accuracy_score(y_test, y_pred)
        precision = precision_score(y_test, y_pred)
        recall = recall_score(y_test, y_pred)
        fmeasure = fmeasure_score(precision, recall)
        specificity = specificity_score(conf_matrix)
        mcc = matthews_corrcoef(y_test, y_pred)
        npv = npv_score(conf_matrix)
        fpr = fpr_score(conf_matrix)
        fnr = fnr_score(conf_matrix)

        metrics.append({
            'Model': name,
            'Split': split_name,
            'Accuracy': accuracy,
            'Precision': precision,

```

```

        'Recall': recall,
        'Fmeasure': fmeasure,
        'Specificity': specificity,
        'MCC': mcc,
        'NPV': npv,
        'FPR': fpr,
        'FNR': fnr
    })
    return metrics

# Example model definitions
models1 = [ ('RandomForest', RandomForestClassifier()), ('LogisticRegression', Log
models2 = [ ('RandomForest', RandomForestClassifier()), ('LogisticRegression', Log

# Example train/test splits
X_train1, X_test1, y_train1, y_test1 = X_train1, X_test1, y_train1, y_test1
X_train2, X_test2, y_train2, y_test2 = X_train2, X_test2, y_train2, y_test2



# Collect all metrics for both splits
metrics_split1 = get_all_metrics(models1, X_train1, X_test1, y_train1, y_test1,
metrics_split2 = get_all_metrics(models2, X_train2, X_test2, y_train2, y_test2,



# Combine results into a DataFrame
metrics_df = pd.DataFrame(metrics_split1 + metrics_split2)


# Function to plot metrics
def plot_metrics(metrics_df, metric_name):
    plt.figure(figsize=(7,5))
    sns.barplot(x='Split', y=metric_name, hue='Model', data=metrics_df, palette=
plt.title(f'{metric_name} of Different Models at Various Data Splits')
plt.xlabel('Data Split')
plt.ylabel(metric_name)
plt.ylim(0, 1)
plt.legend(title='Model', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()



# Plot all metrics
for metric in ['Accuracy', 'Precision', 'Recall', 'Fmeasure', 'Specificity', 'MC
    plot_metrics(metrics_df, metric)


```











```
In [ ]:
In [ ]:
In [ ]:
In [ ]:
```